

model_project_eeg

January 16, 2023

1 EEG Project: Model Answer

1.1 Import Functions

```
[1]: from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier,□
      ↪GradientBoostingClassifier, AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC, LinearSVC
from sklearn.neural_network import MLPClassifier
from sklearn.neighbors import KNeighborsClassifier

from sklearn.preprocessing import Normalizer, StandardScaler
from sklearn.model_selection import train_test_split, StratifiedShuffleSplit,□
      ↪GridSearchCV, permutation_test_score
from sklearn.metrics import roc_curve, roc_auc_score, auc
from sklearn.decomposition import PCA
from sklearn.mixture import GaussianMixture

from matplotlib.pyplot import subplots

from numpy import arange, linspace, interp, asarray, mean, var, sqrt
from numpy import cumsum, equal, count_nonzero, invert, where, std
from numpy import float32, fill_diagonal, corrcoef

from pandas import read_csv
```

2 Data Handling

2.1 Import Data and Generic Checks

- Import data to dataframe
- Print the data types to find categorical columns
- Print all column labels
- Print the head of the dataframe

```
[2]: RANDOM_STATE = 1234
DATA_PATH = 'data/eeg_eye_state.csv'
```

```

df = read_csv(DATA_PATH)

print(df.dtypes)
print('')

columns = df.columns

print(columns)
print('')

df.head()

```

```

AF3           float64
F7            float64
F3            float64
FC5           float64
T7            float64
P7            float64
O1            float64
O2            float64
P8            float64
T8            float64
FC6           float64
F4            float64
F8            float64
AF4           float64
eyeDetection    bool
dtype: object

```

```

Index(['AF3', 'F7', 'F3', 'FC5', 'T7', 'P7', 'O1', 'O2', 'P8', 'T8', 'FC6',
       'F4', 'F8', 'AF4', 'eyeDetection'],
      dtype='object')

```

	AF3	F7	F3	FC5	T7	P7	O1	O2	\
0	4329.23	4009.23	4289.23	4148.21	4350.26	4586.15	4096.92	4641.03	
1	4324.62	4004.62	4293.85	4148.72	4342.05	4586.67	4097.44	4638.97	
2	4327.69	4006.67	4295.38	4156.41	4336.92	4583.59	4096.92	4630.26	
3	4328.72	4011.79	4296.41	4155.90	4343.59	4582.56	4097.44	4630.77	
4	4326.15	4011.79	4292.31	4151.28	4347.69	4586.67	4095.90	4627.69	
	P8	T8	FC6	F4	F8	AF4	eyeDetection		
0	4222.05	4238.46	4211.28	4280.51	4635.90	4393.85		False	
1	4210.77	4226.67	4207.69	4279.49	4632.82	4384.10		False	
2	4207.69	4222.05	4206.67	4282.05	4628.72	4389.23		False	
3	4217.44	4235.38	4210.77	4287.69	4632.31	4396.41		False	
4	4210.77	4244.10	4212.82	4288.21	4632.82	4398.46		False	

2.2 Specific Checks

- Check for missig values
- Drop rows (if any) that contain missing values

```
[3]: has_missing_values = df.isnull().any()

print(has_missing_values)
```

```
AF3      False
F7       False
F3       False
FC5      False
T7       False
P7       False
O1       False
O2       False
P8       False
T8       False
FC6      False
F4       False
F8       False
AF4      False
eyeDetection  False
dtype: bool
```

```
[4]: df = df.dropna(axis=0)
```

```
[5]: label_counts = df['eyeDetection'].value_counts()

n_samples = label_counts.sum()

print('Total number of samples:', n_samples)

print(label_counts)

print('')

percentage_count = label_counts * 100 / n_samples

print('Fraction (%): ')
print(percentage_count.round(2))
```

```
Total number of samples: 14980
False    8257
True     6723
Name: eyeDetection, dtype: int64
```

```
Fraction (%):
```

```
False    55.12
True     44.88
Name: eyeDetection, dtype: float64
```

2.3 Numpy array

- Convert data to numpy array
- Print shape of numpy array

```
[6]: data = df.to_numpy()
```

```
data = data[:, :14]

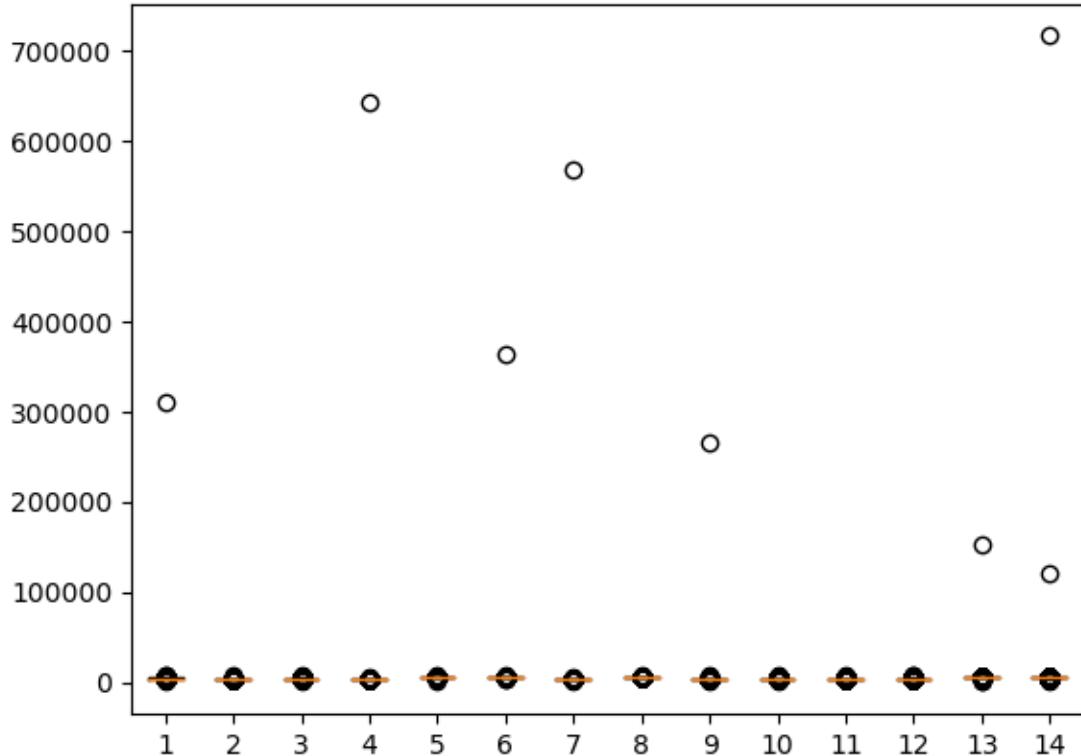
print('Data shape: ', data.shape)
print('')
```

```
Data shape: (14980, 14)
```

```
## Make Boxplot
```

```
[7]: fig, ax = subplots()

ax.boxplot(data);
```



2.4 Remove Outliers and redo Boxplot

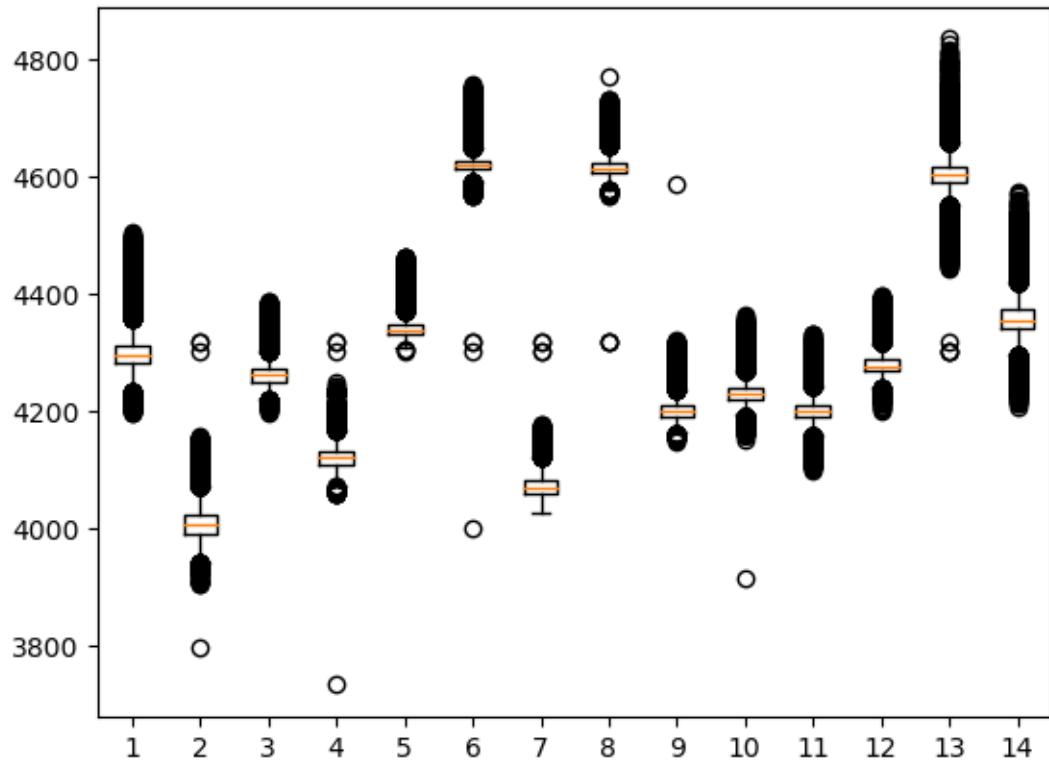
Comment: All values are converted to float because integers and boolean values will create problems.

```
[8]: data[data > 5000] = mean(data)
data[data < 3600] = mean(data)

data = data.astype(float32)

fig, ax = subplots()

ax.boxplot(data);
```

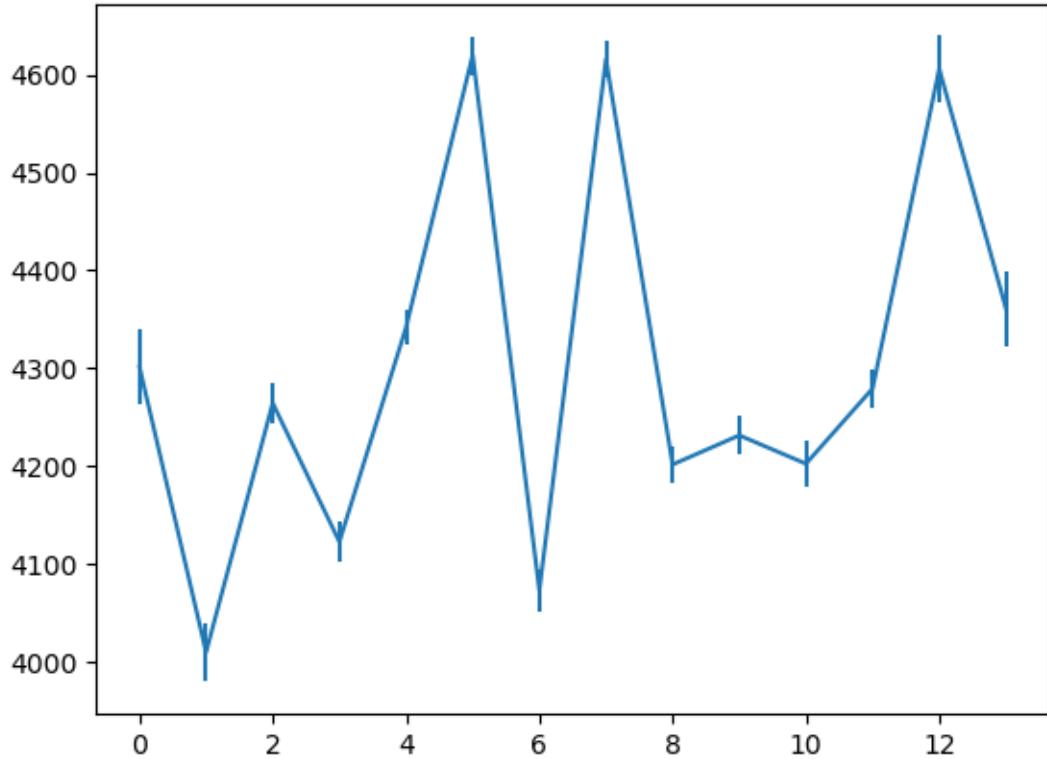


2.5 Means and standard deviations

```
[9]: df_means = mean(data, axis=0)
df_stds = std(data, axis=0)

fig, ax = subplots()

ax.errorbar(range(len(df_means)), df_means, yerr=df_stds);
```



2.6 Correlation Matrix

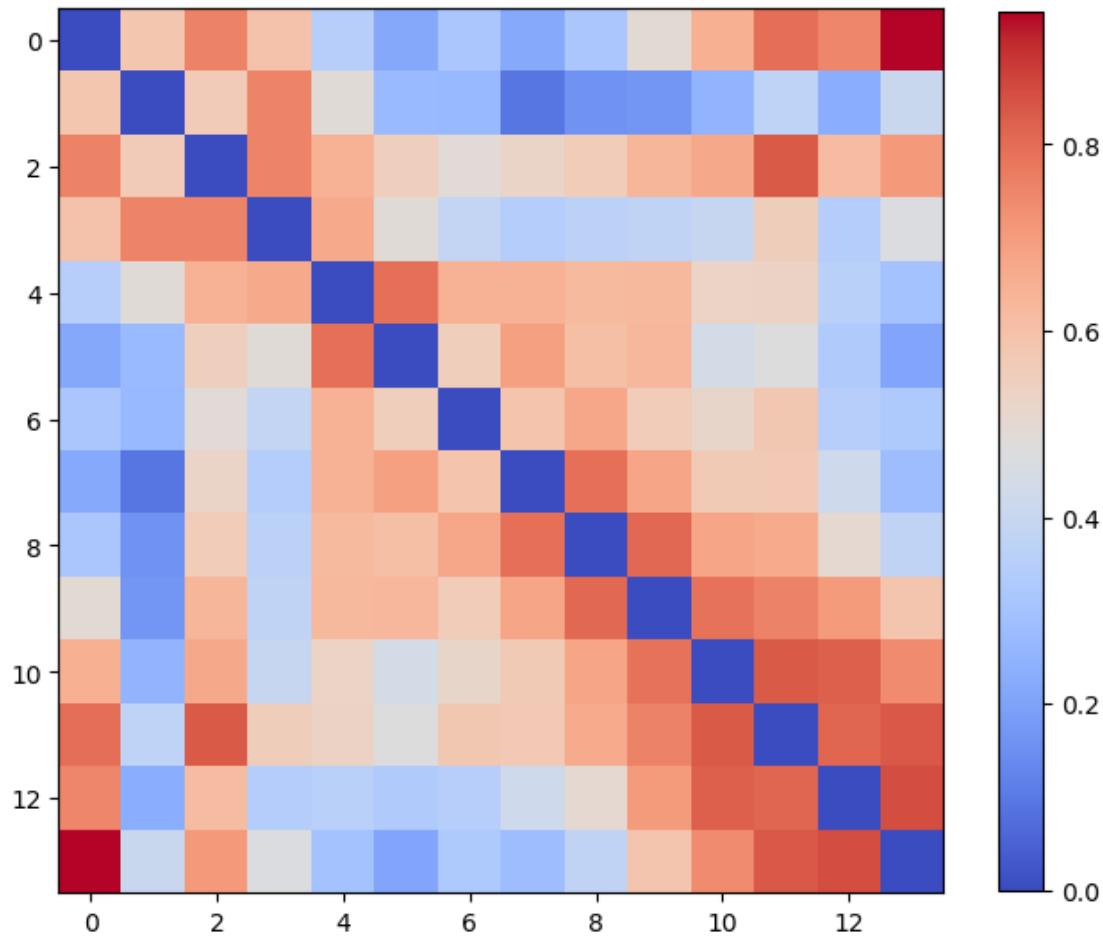
```
[10]: data_corr = corrcoef(data, rowvar=False)

fill_diagonal(data_corr, 0)

fig, ax = subplots(figsize = (8,8))

im = ax.imshow(data_corr, cmap='coolwarm');

fig.colorbar(im, orientation='vertical', shrink=0.8);
```



3 Dimensionality Reduction

- Use PCA on all components
- Plot explained variance as a function of component
- Calculate number of components required to keep 99% variance
- Create labels array
- Scatter plot of main components

```
[11]: nComp = data.shape[1] # Number of PCs to be returned
#trainIdx = np.random.binomial(1,0.9,size=filectr)

PCA_inst = PCA(n_components=nComp, whiten=True)

data_PCA = PCA_inst.fit(data)

cumExpVar = cumsum(data_PCA.explained_variance_ratio_)
```

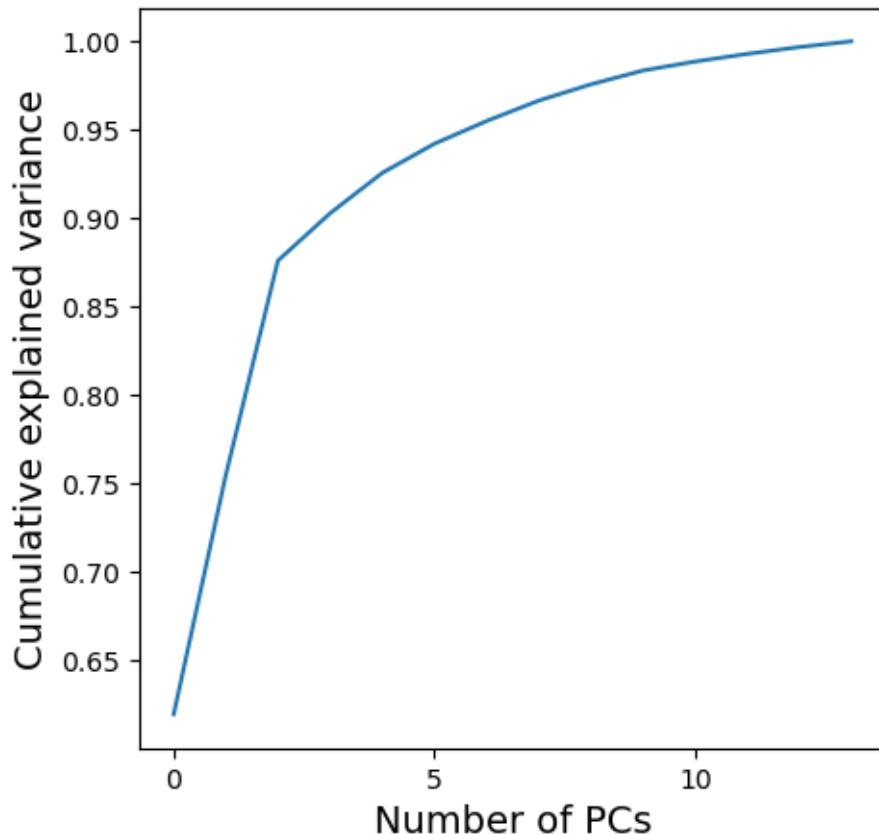
```

fig, ax = subplots(figsize=(5, 5))

im = ax.plot( range(nComp), cumExpVar )

ax.set_xticks(arange(0,nComp,5));
ax.set_xlabel( 'Number of PCs', fontsize=14);
ax.set_ylabel('Cumulative explained variance', fontsize=14);

```



```

[12]: threshold = 0.99

keepPC = [pc for pc in range(nComp) if cumExpVar[pc]>=threshold][0]

print('Number of features needed to explain {:.2f} fraction of total variance is {:d}. '.format(threshold, keepPC))

```

Number of features needed to explain 0.99 fraction of total variance is 11.

```

[13]: labels = asarray(df['eyeDetection'])
y = labels.astype(int)

```

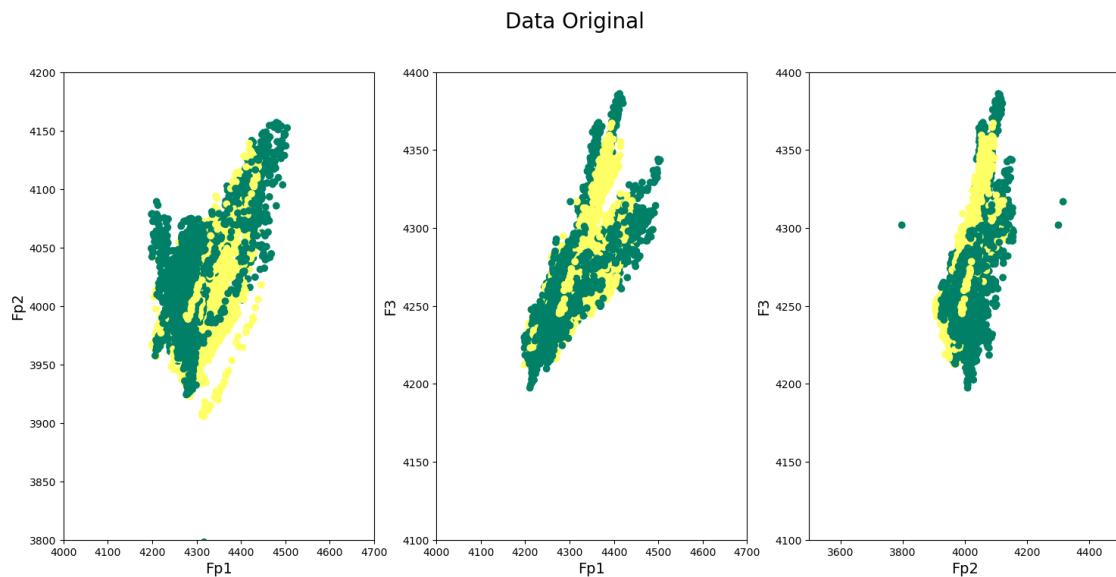
```
[14]: fig, ax = subplots(ncols=3, figsize=(18, 8))

fig.suptitle('Data Original', fontsize=20)

ax[0].scatter(data[:,0], data[:,1], c=y, cmap='summer');
ax[0].set_xlabel('Fp1', fontsize=14); ax[0].set_xlim([4000,4700]);
ax[0].set_ylabel('Fp2', fontsize=14); ax[0].set_ylim([3800,4200]);

ax[1].scatter(data[:,0], data[:,2], c=y, cmap='summer');
ax[1].set_xlabel('Fp1', fontsize=14); ax[1].set_xlim([4000,4700]);
ax[1].set_ylabel('F3', fontsize=14); ax[1].set_ylim([4100,4400]);

ax[2].scatter(data[:,1], data[:,2], c=y, cmap='summer');
ax[2].set_xlabel('Fp2', fontsize=14); ax[2].set_xlim([3500,4500]);
ax[2].set_ylabel('F3', fontsize=14); ax[2].set_ylim([4100,4400]);
```



```
[15]: new_data = PCA_inst.transform(data)[:,range(keepPC)]

fig, ax = subplots(ncols=3, figsize=(18, 8))

fig.suptitle('Data in new Feature space', fontsize=20)

ax[0].scatter(new_data[:,0], new_data[:,1], c=y, cmap='summer');
ax[0].set_xlabel('PC1', fontsize=14); ax[0].set_xlim([-5,5]);
ax[0].set_ylabel('PC2', fontsize=14); ax[0].set_ylim([-5,5]);

ax[1].scatter(new_data[:,0], new_data[:,2], c=y, cmap='summer');
ax[1].set_xlabel('PC1', fontsize=14); ax[1].set_xlim([-5,5]);
```

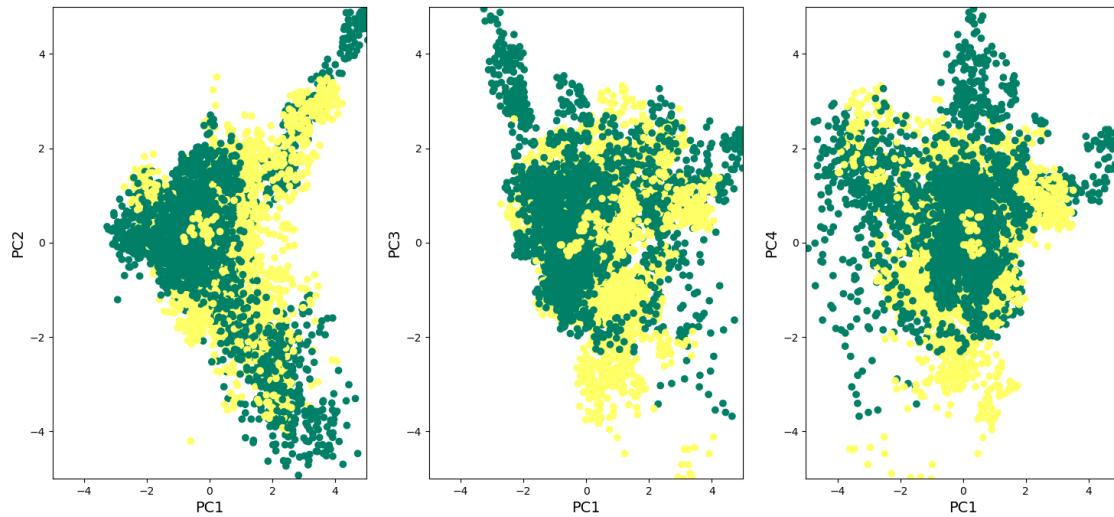
```

ax[1].set_ylabel('PC3', fontsize=14); ax[1].set_ylim([-5,5]);

ax[2].scatter(new_data[:,1], new_data[:,2], c=y, cmap='summer');
ax[2].set_xlabel('PC1', fontsize=14); ax[2].set_xlim([-5,5]);
ax[2].set_ylabel('PC4', fontsize=14); ax[2].set_ylim([-5,5]);

```

Data in new Feature space



3.1 Scaling

```

[16]: scaler = StandardScaler()

data_scaled = scaler.fit_transform(data)

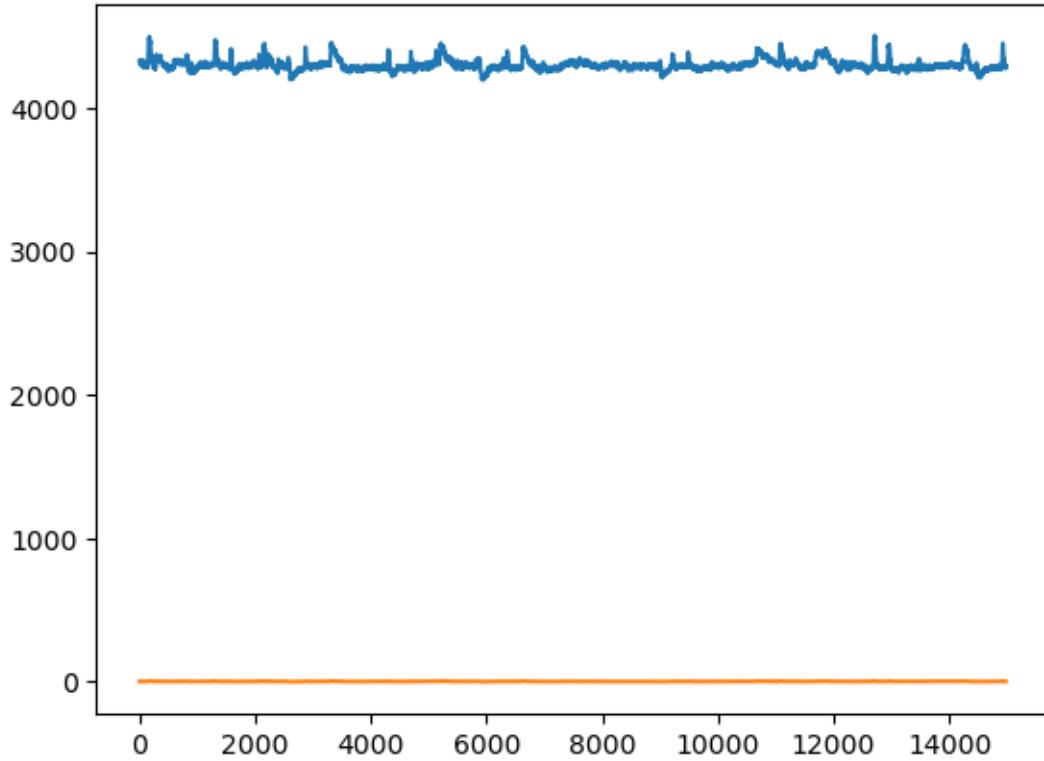
print(data.shape, data_scaled.shape)

fig, ax = subplots()

ax.plot(data[:, 0])
ax.plot(data_scaled[:, 0]);

```

(14980, 14) (14980, 14)



4 Clustering

```
[17]: X = data_scaled  
  
y = y  
  
clf = GaussianMixture(n_components=2)  
  
clf.fit(X)  
  
y_predict = clf.predict(X)  
  
  
from sklearn.metrics.cluster import adjusted_rand_score  
  
scoring = adjusted_rand_score(y, y_predict)  
  
print(scoring)
```

0.012647818008879305

The adjusted rand score is around zero. The GMM cannot separate the two classes.

5 Supervised Learning

5.1 Classification of Original Data

```
[18]: from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier,  
      GradientBoostingClassifier, AdaBoostClassifier  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.svm import SVC, LinearSVC  
from sklearn.neural_network import MLPClassifier  
from sklearn.neighbors import KNeighborsClassifier  
  
classifiers = {  
    'Random Forest': RandomForestClassifier(random_state=RANDOM_STATE),  
    'AdaBoost (Random Forest)':  
        AdaBoostClassifier(RandomForestClassifier(random_state=RANDOM_STATE)),  
    'Extra Trees': ExtraTreesClassifier(random_state=RANDOM_STATE),  
    'AdaBoost (Extra Tree)':  
        AdaBoostClassifier(ExtraTreesClassifier(random_state=RANDOM_STATE)),  
    'Decision Tree': DecisionTreeClassifier(random_state=RANDOM_STATE),  
    'SVC (RBF)': SVC(random_state=RANDOM_STATE),  
    'SVC (Linear)': LinearSVC(random_state=RANDOM_STATE),  
    'Multi-layer Perceptron': MLPClassifier(max_iter=5000,  
        random_state=RANDOM_STATE)  
}
```

```
[19]: RANDOM_STATE = 1234  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.3,  
                                                random_state=RANDOM_STATE, shuffle=True)  
  
scores = list()  
name_list = list()  
  
for name, clf in classifiers.items():  
    # Training the model using training data:  
    clf.fit(X_train, y_train)  
  
    y_pred = clf.predict(X_test)  
  
    # Evaluating the score using test data:  
    score = clf.score(X_test, y_test)  
    scores.append(score)  
    name_list.append(name)  
    print(name, score)
```

```
Random Forest 0.9272363150867824  
AdaBoost (Random Forest) 0.92456608811749  
Extra Trees 0.9421450823319982
```

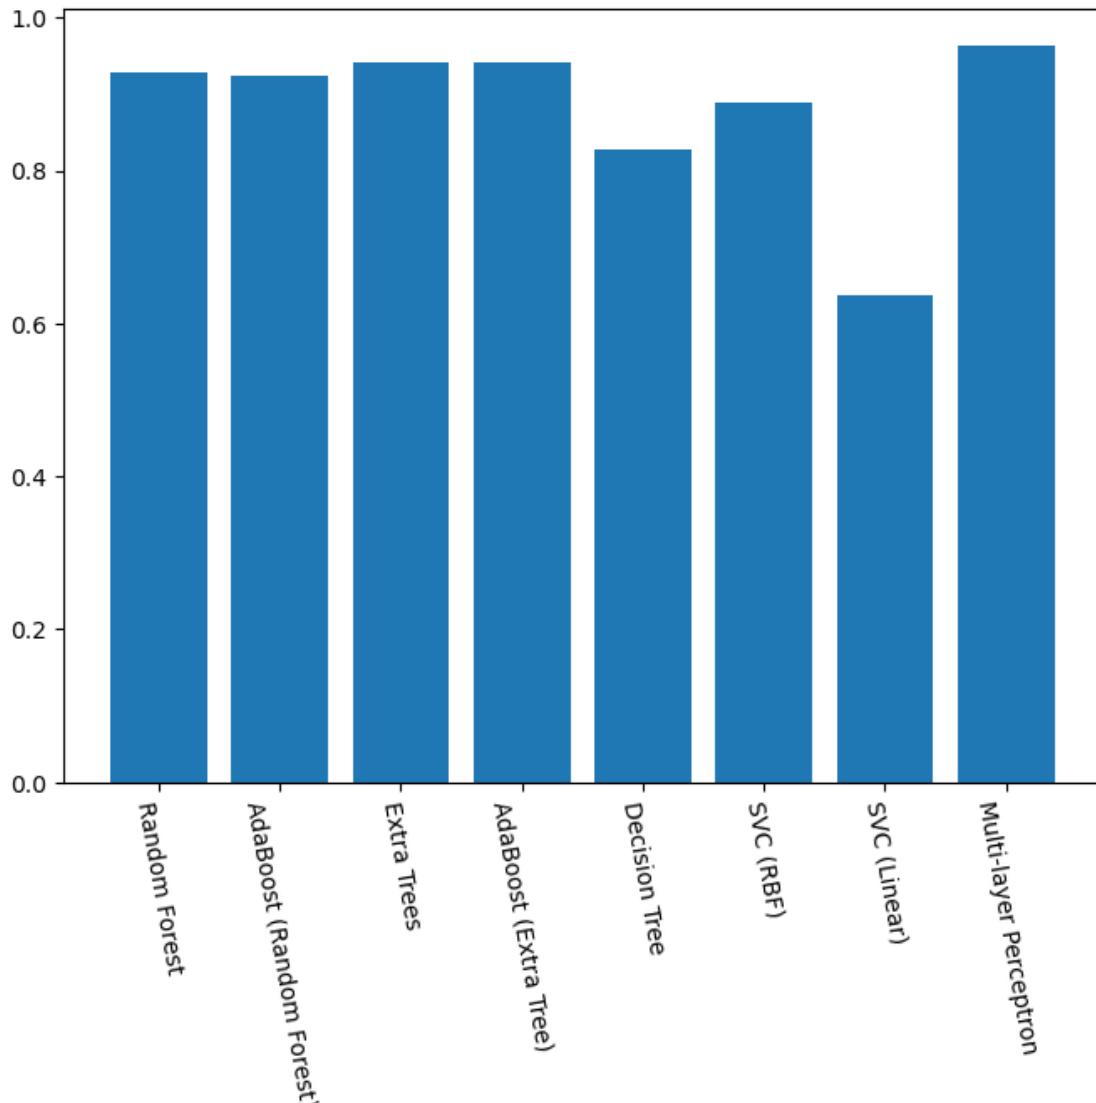
```
AdaBoost (Extra Tree) 0.9417000445037829
Decision Tree 0.8266577659101023
SVC (RBF) 0.8887405429461505

/Users/geroldbaier/opt/anaconda3/lib/python3.8/site-
packages/sklearn/svm/_base.py:1206: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.

    warnings.warn(
SVC (Linear) 0.6372941700044504
Multi-layer Perceptron 0.9639519359145527
```

```
[20]: fig, ax = subplots(figsize=(8,6))
```

```
bins = arange(8)+0
ax.bar(arange(len(scores)), scores)
ax.set_xticks(bins)
ax.set_xticklabels(name_list, rotation=-80);
```



The ensemble classifiers achieve a good classification. So does the artificial neural network (MLP).

```
[21]: from sklearn.metrics import roc_curve, roc_auc_score

fig, all_axes = subplots(figsize=[15, 10], ncols=4, nrows=2, sharey=True, sharex=True)

for ax, (name, clf) in zip(all_axes.ravel(), classifiers.items()):
    clf.fit(X_train, y_train)

    # Checking whether or not the object has `decision_function`:
    if hasattr(clf, 'decision_function'):
        # If it does:
```

```

        y_score = clf.decision_function(X_test)
    else:
        # Otherwise:
        y_score = clf.predict_proba(X_test)[:, 1] # only one probability is
        ↪needed

    # Obtaining the x- and y-axis values for the ROC curve:
    fpr, tpr, thresh = roc_curve(y_test, y_score)

    # Obtaining the AUC value:
    roc_auc = roc_auc_score(y_test, y_score)

    ax.plot(fpr, tpr, lw=2)
    ax.plot([0, 1], [0, 1], lw=1, linestyle='--')

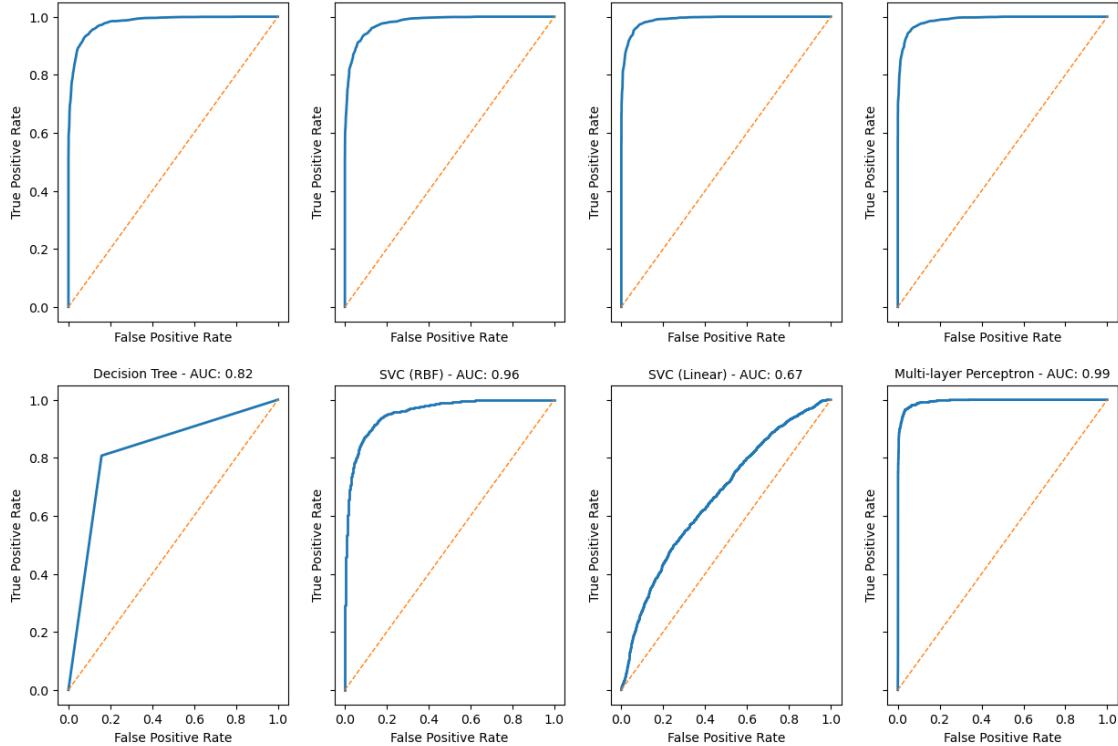
    ax.set_xlabel('False Positive Rate')
    ax.set_ylabel('True Positive Rate')

    label = '{} - AUC: {:.2f}'.format(name, roc_auc)
    ax.set_title(label, fontsize=10)

```

/Users/geroldbaier/opt/anaconda3/lib/python3.8/site-packages/sklearn/svm/_base.py:1206: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn(



5.2 Classification of Principal Components

```
[22]: RANDOM_STATE = 1234

X_PCA = new_data

X_train_pca, X_test_pca, y_train_pca, y_test_pca = train_test_split(X_PCA, y, ↴
    test_size=.3, random_state=RANDOM_STATE, shuffle=True)

new_scores = list()
name_list = list()

for name, clf in classifiers.items():
    # Training the model using training data:
    clf.fit(X_train, y_train)

    y_pred = clf.predict(X_test)

    # Evaluating the score using test data:
    score = clf.score(X_test, y_test)
    new_scores.append(score)
    name_list.append(name)
    print(name, score)

print('Complete')
```

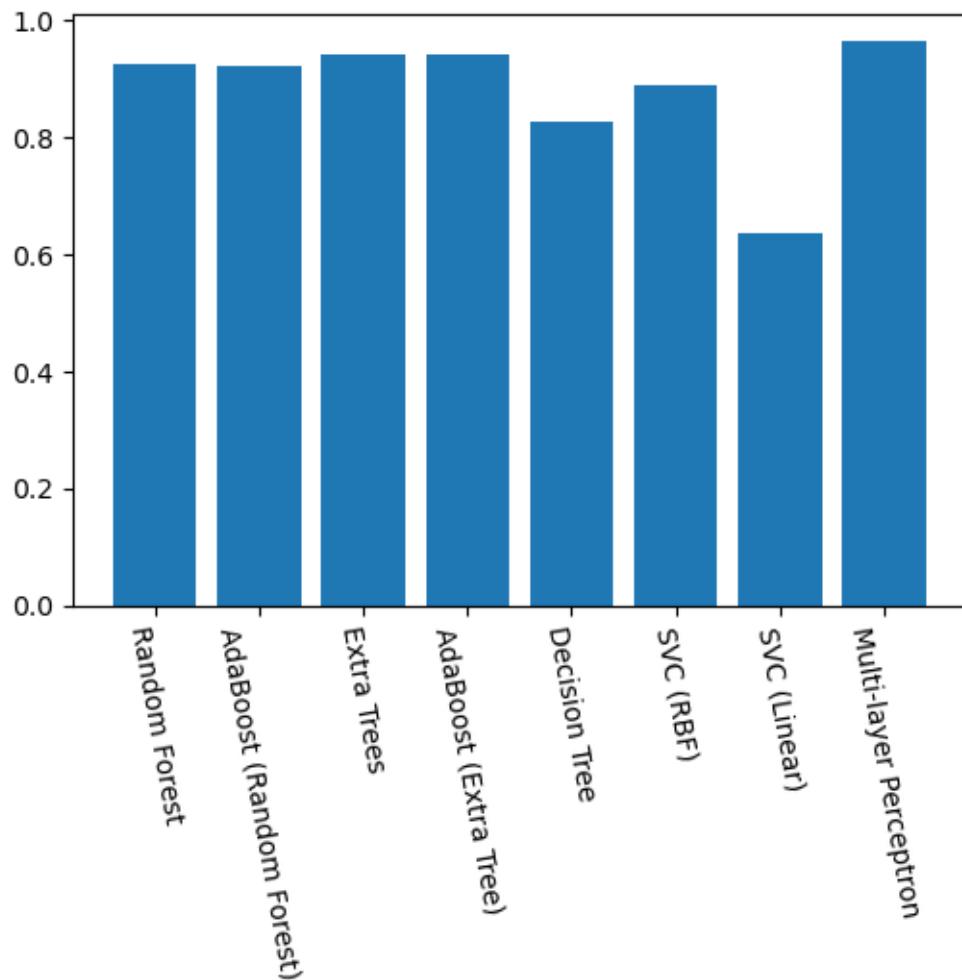
```
Random Forest 0.9272363150867824
AdaBoost (Random Forest) 0.9218958611481975
Extra Trees 0.9421450823319982
AdaBoost (Extra Tree) 0.9419225634178905
Decision Tree 0.8266577659101023
SVC (RBF) 0.8887405429461505

/Users/geroldbaier/opt/anaconda3/lib/python3.8/site-
packages/sklearn/svm/_base.py:1206: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
    warnings.warn(
SVC (Linear) 0.6372941700044504
Multi-layer Perceptron 0.9639519359145527
Complete
```

```
[26]: fig, ax = subplots(figsize=(6,4))

bins = arange(8)+0
ax.bar(arange(len(new_scores)), new_scores)
```

```
ax.set_xticks(bins)
ax.set_xticklabels(name_list, rotation=-80);
```



The prediction from the reduced data is similar to that of the original data. That will help to speed up the analysis of large datasets.

```
[27]: from sklearn.metrics import roc_curve, roc_auc_score

fig, all_axes = subplots(figsize=[15, 10], ncols=4, nrows=2, sharey=True, sharex=True)

for ax, (name, clf) in zip(all_axes.ravel(), classifiers.items()):
    clf.fit(X_train_pca, y_train_pca)

    # Checking whether or not the object has `decision_function`:
    if hasattr(clf, 'decision_function'):
```

```

# If it does:
y_score_pca = clf.decision_function(X_test_pca)
else:
    # Otherwise:
    y_score_pca = clf.predict_proba(X_test_pca)[:, 1] # only one probability
    ↪is needed

# Obtaining the x- and y-axis values for the ROC curve:
fpr_pca, tpr_pca, thresh = roc_curve(y_test_pca, y_score_pca)

# Obtaining the AUC value:
roc_auc_pca = roc_auc_score(y_test_pca, y_score_pca)

ax.plot(fpr_pca, tpr_pca, lw=2)
ax.plot([0, 1], [0, 1], lw=1, linestyle='--')

ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')

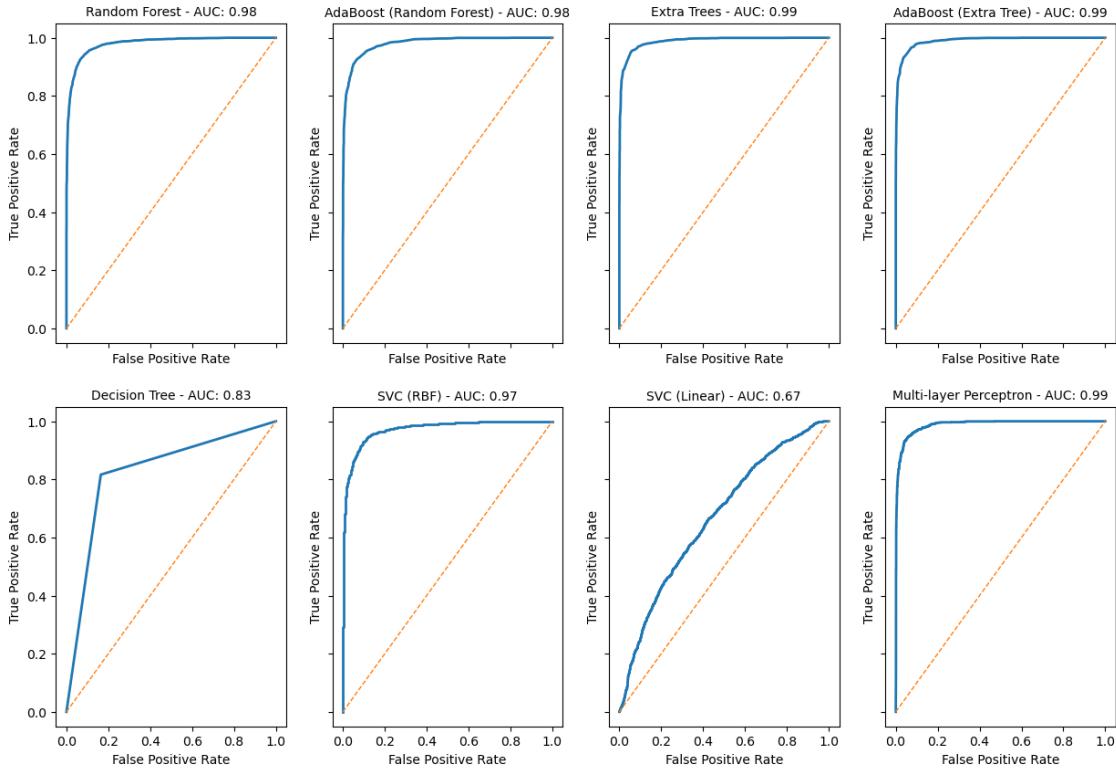
label = '{} - AUC: {:.2f}'.format(name, roc_auc_pca)
ax.set_title(label, fontsize=10)

```

```

/Users/geroldbaier/opt/anaconda3/lib/python3.8/site-
packages/sklearn/svm/_base.py:1206: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(

```



The reduced dataset achieves similar scores while reducing time of computation.

5.3 Feature Importances

```
[28]: clf_data = ExtraTreesClassifier(random_state=RANDOM_STATE)

clf_data.fit(X, y)

importances = clf_data.feature_importances_

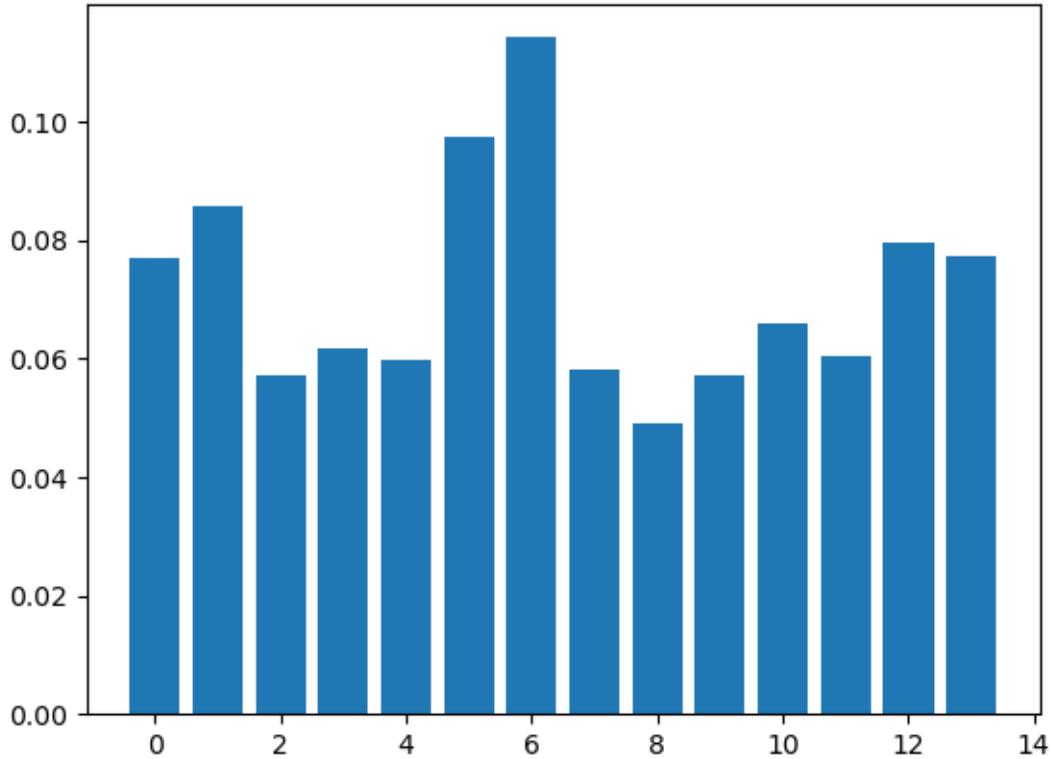
print('Relative feature importance:')

bins = arange(importances.shape[0])

fig, ax = subplots()

ax.bar(bins, importances);
```

Relative feature importance:



```
[29]: ind_max = where(importances == max(importances))

channels = df.columns

print('The name of the most important channel is: ', channels[int(ind_max[0])])
```

The name of the most important channel is: 01

[]:

```
[30]: clf_data_dectree = DecisionTreeClassifier(random_state=RANDOM_STATE)

clf_data_dectree.fit(X, y)

importances = clf_data_dectree.feature_importances_

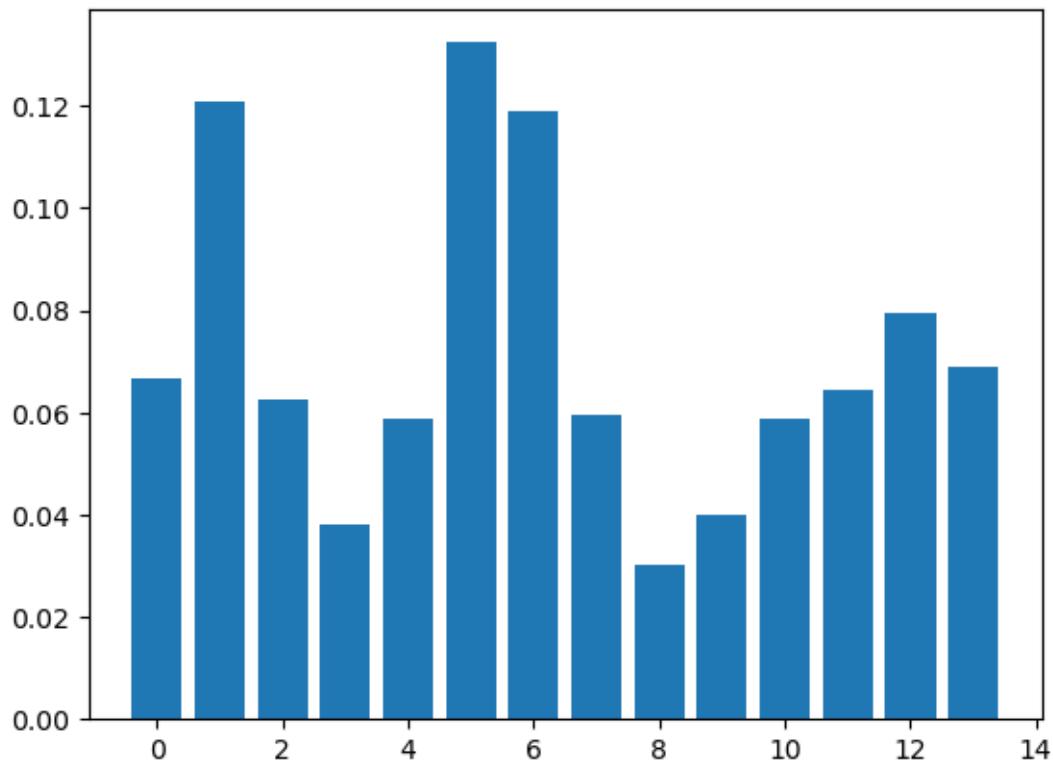
print('Relative feature importance:')

bins = arange(importances.shape[0])

fig, ax = subplots()
```

```
ax.bar(bins, importances);
```

Relative feature importance:



```
[31]: ind_max = where(importances == max(importances))

channels = df.columns

print('The name of the most important channel is: ', channels[int(ind_max[0])])
```

The name of the most important channel is: P7